

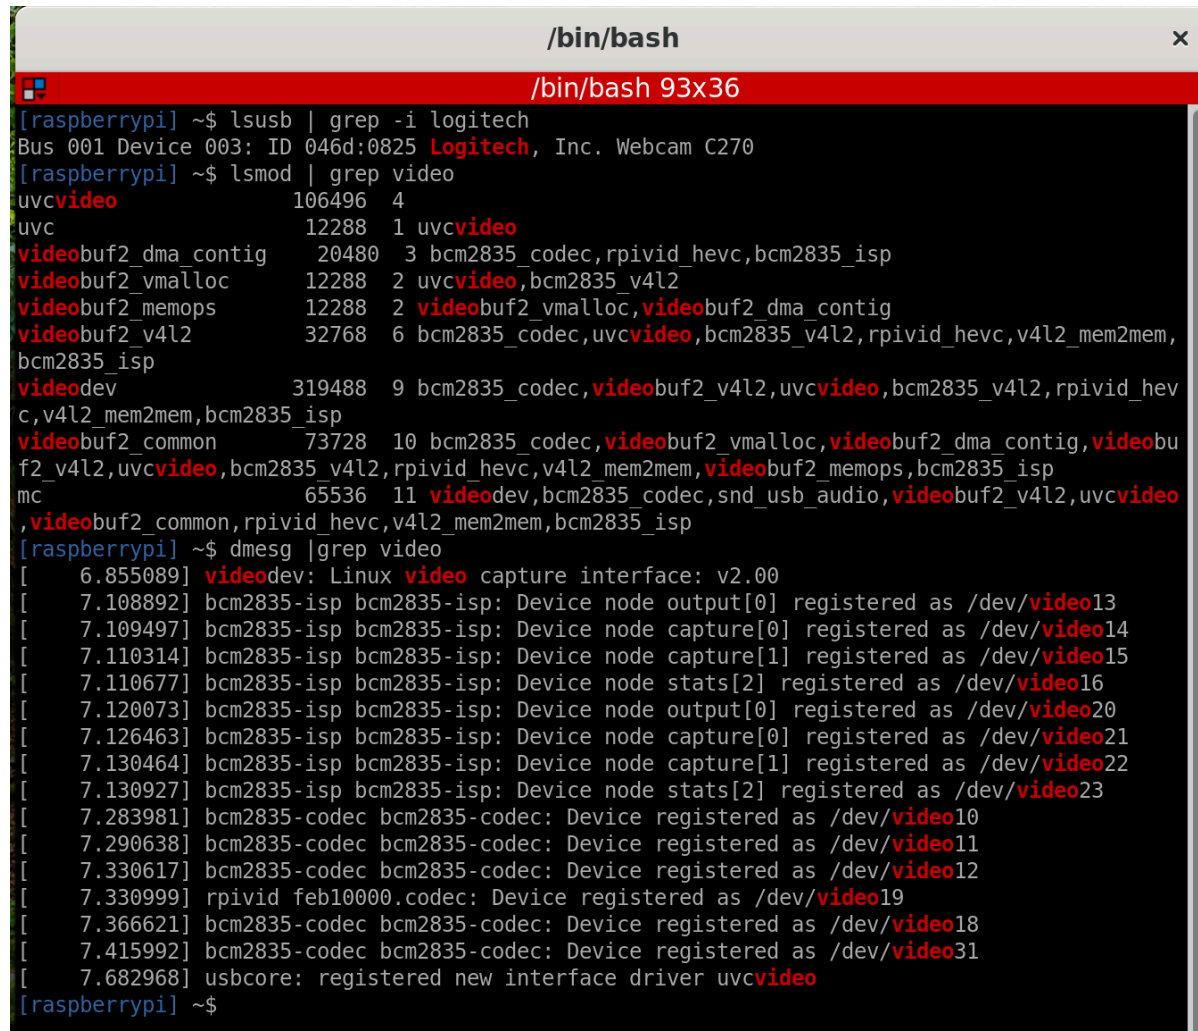
# ECEE 5623: Real-Time Systems

## Exercise 4: Real-Time Continuous Media

Due: July 12th, 2024

Jack Center

### Part 1: Verify Camera



```
/bin/bash
/bin/bash 93x36
[raspberry] ~$ lsusb | grep -i logitech
Bus 001 Device 003: ID 046d:0825 Logitech, Inc. Webcam C270
[raspberry] ~$ lsmod | grep video
uvcvideo                106496  4
uvc                      12288  1 uvcvideo
videobuf2_dma_contig    20480  3 bcm2835_codec,rpivid_hevc,bcm2835_isp
videobuf2_vmalloc       12288  2 uvcvideo,bcm2835_v4l2
videobuf2_memops        12288  2 videobuf2_vmalloc,videobuf2_dma_contig
videobuf2_v4l2          32768  6 bcm2835_codec,uvcvideo,bcm2835_v4l2,rpivid_hevc,v4l2_mem2mem,
bcm2835_isp
videodev                 319488  9 bcm2835_codec,videobuf2_v4l2,uvcvideo,bcm2835_v4l2,rpivid_hevc,
v4l2_mem2mem,bcm2835_isp
videobuf2_common        73728  10 bcm2835_codec,videobuf2_vmalloc,videobuf2_dma_contig,videobu
f2_v4l2,uvcvideo,bcm2835_v4l2,rpivid_hevc,v4l2_mem2mem,videobuf2_memops,bcm2835_isp
mc                       65536  11 videodev,bcm2835_codec,snd_usb_audio,videobuf2_v4l2,uvcvideo
,videobuf2_common,rpivid_hevc,v4l2_mem2mem,bcm2835_isp
[raspberry] ~$ dmesg |grep video
[ 6.855089] videodev: Linux video capture interface: v2.00
[ 7.108892] bcm2835-isp bcm2835-isp: Device node output[0] registered as /dev/video13
[ 7.109497] bcm2835-isp bcm2835-isp: Device node capture[0] registered as /dev/video14
[ 7.110314] bcm2835-isp bcm2835-isp: Device node capture[1] registered as /dev/video15
[ 7.110677] bcm2835-isp bcm2835-isp: Device node stats[2] registered as /dev/video16
[ 7.120073] bcm2835-isp bcm2835-isp: Device node output[0] registered as /dev/video20
[ 7.126463] bcm2835-isp bcm2835-isp: Device node capture[0] registered as /dev/video21
[ 7.130464] bcm2835-isp bcm2835-isp: Device node capture[1] registered as /dev/video22
[ 7.130927] bcm2835-isp bcm2835-isp: Device node stats[2] registered as /dev/video23
[ 7.283981] bcm2835-codec bcm2835-codec: Device registered as /dev/video10
[ 7.290638] bcm2835-codec bcm2835-codec: Device registered as /dev/video11
[ 7.330617] bcm2835-codec bcm2835-codec: Device registered as /dev/video12
[ 7.330999] rpivid febl10000.codec: Device registered as /dev/video19
[ 7.366621] bcm2835-codec bcm2835-codec: Device registered as /dev/video18
[ 7.415992] bcm2835-codec bcm2835-codec: Device registered as /dev/video31
[ 7.682968] usbcore: registered new interface driver uvcvideo
[raspberry] ~$
```

Figure 1: Output from system commands for part 1.

- `lsusb` lists USB devices and shows the Logitech camera.
- `lsmod` shows the status of modules and shows the UVC driver is loaded, the size of the module in bytes, and what is using the module.
- `dmesg` shows the kernel ring buffer and shows messages stored in it. It shows the Linux video capture interface is in the buffer.

## Part 2: Install Cheese

```
/bin/bash
/bin/bash 100x21
Start-Date: 2024-07-07 11:45:25
Commandline: apt install cheese
Requested-By: rpi (1000)
Install: libcogl-path20:arm64 (1.22.8-3+b1, automatic), libcanberra-gtk3-0:arm64 (0.30-10, automatic), libclutter-1.0-common:arm64 (1.26.4+dfsg-4, automatic), libcanberra-gtk3-module:arm64 (0.30-10, automatic), fonts-dejavu:arm64 (2.37-6, automatic), yelp-xsl:arm64 (42.1-2, automatic), gstreamer1.0-clutter-3.0:arm64 (3.0.27-3, automatic), libcogl20:arm64 (1.22.8-3+b1, automatic), libclutter-gtk-1.0-0:arm64 (1.8.4-4+b1, automatic), libcheese-gtk25:arm64 (43.0-1, automatic), docbook-xml:arm64 (4.5-12, automatic), sgml-data:arm64 (2.0.11+nmu1, automatic), libcogl-common:arm64 (1.22.8-3, automatic), cheese:arm64 (43.0-1), sound-theme-freedesktop:arm64 (0.8-2, automatic), gnome-video-effects:arm64 (0.5.0-1, automatic), libyelp0:arm64 (42.2-1, automatic), yelp:arm64 (42.2-1, automatic), libcogl-pango20:arm64 (1.22.8-3+b1, automatic), libcheese8:arm64 (43.0-1, automatic), fonts-dejavu-extra:arm64 (2.37-6, automatic), libclutter-1.0-0:arm64 (1.26.4+dfsg-4, automatic), cheese-common:arm64 (43.0-1, automatic), libclutter-gst-3.0-0:arm64 (3.0.27-3, automatic), libcanberra0:arm64 (0.30-10, automatic)
End-Date: 2024-07-07 11:45:42
~
~
~
18,30 Bot
```

Figure 2.1: Summary log for installation of Cheese.

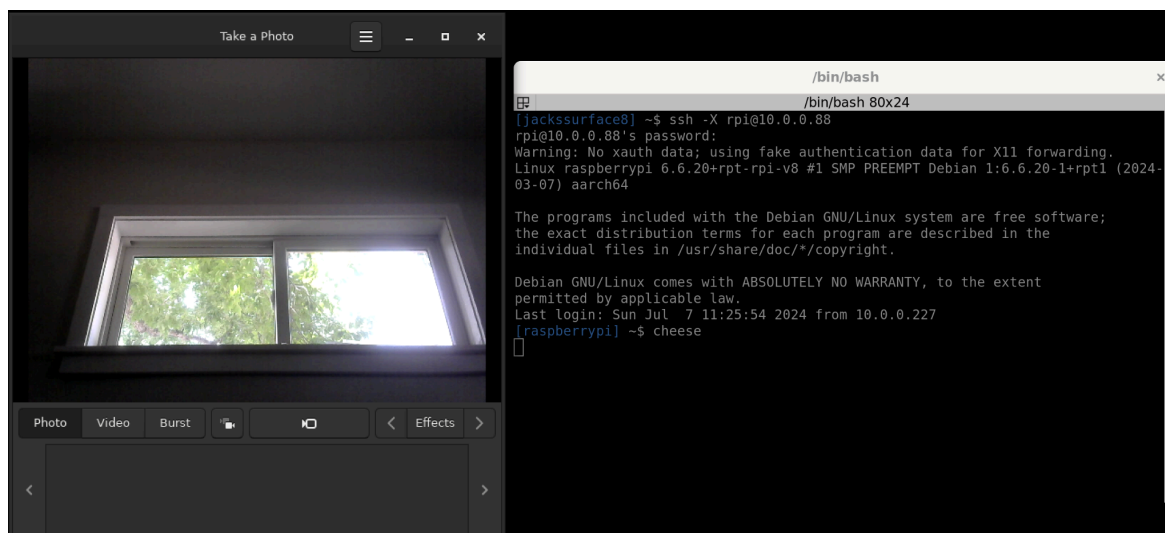
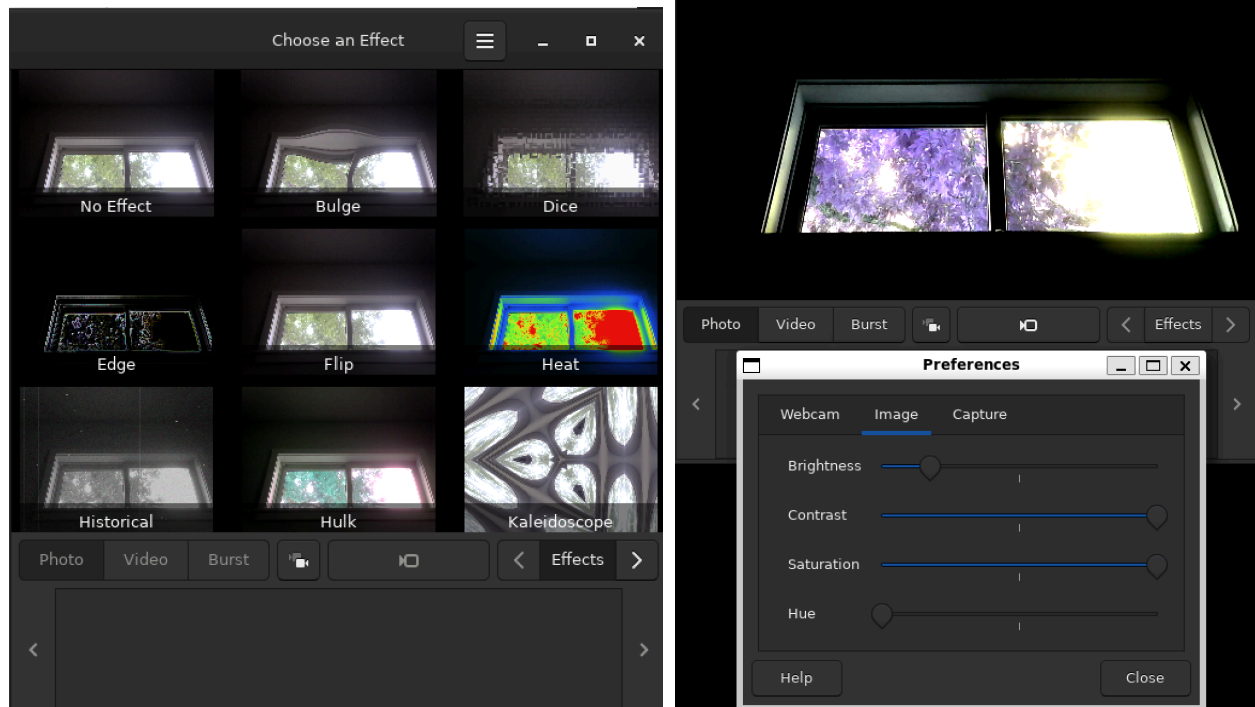


Figure 2.2: GUI output for Cheese.



*Figure 2.3: Screenshot of effects tools and an example using it.*

### Part 3: Continuous Streaming

Figure 3 shows continuous capture is working. The image files can be seen in the terminal and eom window. Table 3 compares the results of processing 1,800 graymap and color frames. The median time per frame shows that graymap processing is about three times faster than color frame processing. This makes sense since there are three times the data in a color frame. The worst-case execution times varied significantly, which also brought up the mean execution time for the color frame processing. Considering the final project needs a 10Hz processing time, neither of these frame rates is sufficient to meet the Nyquist requirement of a 5Hz signal. This means a different method of capturing and processing the images will need to be implemented to separate the image capture from writing to flash so the correct rate can be achieved.

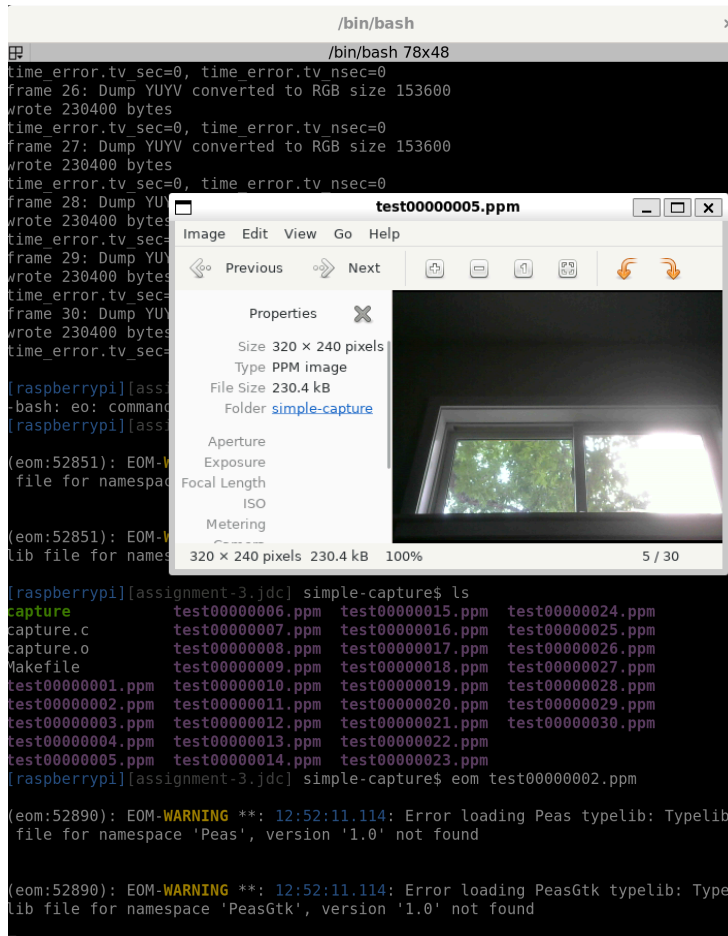


Figure 3: Screenshot of using continuous capture.

	PGM	PPM
Mean Execution Time (ms)	5.437	24.391
Median Execution Time (ms)	5.288	14.052
Worst Case Execution Time (ms)	133.798	3,558.454

Table 3: Results from simple-capture-1800 on graymap (PGM) and color (PPM) frames.

#### Part 4: Continuous Transformation

I implemented conversion to grayscale using each of the single color bands and the balanced equation the GIMP uses. Overall, the balanced image is the highest quality, but the green band also does a good job since a lot of the image was green. The blue image is dark, but, as expected, the blue tape is more visible with this conversion.



*Figure 4: Output from grayscale transformations using red, blue, and green color bands along with a balanced combination of the three based on the conversion used by GIMP.*

The code to complete this transformation was based on the starter code in `simple-capture`. There are three major phases outside of setup and teardown: image acquisition, transformation, and write-back to flash. The code updated for this project was in the `process\_image` function under the color conversion option and mostly dealt with the image transformation. The updated algorithm takes the output from the YUV to RGB conversion and loops through each group of RGB values provided. Each of these groups is combined according to the balancing equation used by GIMP:

$$Y = 0.3R + 0.59G + 0.11B$$

These values, which are one-third the size of the original RGB data, are then written to a PGM file in flash memory.

Transformation such as this can be used in machine vision applications to focus in on relevant data. For example, if a camera is pointed at a stop light, it could be used to mask the incoming colors and only focus on red to identify if the light is red. This can be taken even further by focusing on a specific shade of red to isolate only the specific wavelength coming off of the stoplight. Or, this technique could be used to track an identifying marker on a vehicle which

could then be used to calculate how many pixels it has moved from one frame to the next; this type of information could be used to predict or measure the object's motion. This narrowing of scope for what the image is seeing simplifies follow-on tasks that need to interpret what the image means and reduce the likelihood of noise interfering with the analysis.

The best frame rate for the transformation was measured to be 8.213 milliseconds (Table 5). This value is independent of writing to the flash file system since the time measurements for the three stages were separated for Part 5. The best frame rate time without this conversion was 3.018 milliseconds. This indicates that the grayscale transformation more than doubles the time the transformation takes without the conversion.

## Part 5: Deadline Analysis

The results for the transform performance are summarized in Table 5.a. and 5.b.:

	Mean (ms)	Median (ms)	Worst (ms)	Best (ms)
<b>Acquisition</b>	0.016	0.016	0.034	0.007
<b>Transformation</b>	12.298	12.419	24.349	8.213
<b>Write-back</b>	0.416	0.315	72.985	0.217
<b>Total</b>	12.788	12.815	85.466	8.515

*Table 5.a.: Timing metrics from capturing 1,800 color frames and converting them to grayscale using the balanced implementation described in Part 4.*

	Mean (ms)	Median (ms)	Worst (ms)	Best (ms)
<b>Acquisition</b>	0.018	0.018	0.047	0.011
<b>Transformation</b>	5.677	5.741	9.581	3.018
<b>Write-back</b>	0.653	0.484	127.801	0.314
<b>Total</b>	6.412	6.319	133.733	3.390

*Table 5.b.: Timing metrics for capturing 1,800 color frames. This is unique from Table 5.a. because this data does not execute the grayscale transform.*


Reviewing the performance shows that the transformation is typically the longest part of the procedure; however, the write-back can get backed up and cause substantial delays. These delays are substantial enough to cause missed deadlines. The grayscale transformation also has the effect of approximately doubling the total transformation time. This is likely due to how the current implementation first converts to RGB, then converts to grayscale. This could likely be improved by converting directly to grayscale. Overall, this analysis points to two major performance problems to be aware of in the final project: write-back can cause a significant delay and the transformation implementation details are important for service duration.

Based on the average analysis for the transform frame-only write-back, 15 milliseconds was selected for the deadline; however, initial analysis showed that the FIFO implementation was taking closer to 35 milliseconds on average, so that value was used instead for the 1,800 frames. The results in Figure 5, show significant variability in the time for a single iteration of the service to capture, transform, and write back a frame. The first service reliably takes much longer than the other processes, which indicates a need to make sure everything is ready to go before expecting the deadlines to be hit. After the first 250 frames, the process time returns from approximately 35 milliseconds back to the expected 12 millisecond range. However, there are still significant steps in average processing time in the figure that are currently unexplained.

Solving these steps will be crucial for getting predictable results. There are also some long outliers, but these can be explained by the write-back delays discussed previously.

Overall, this implementation does not predictably meet deadlines primarily due to these significant steps in the service duration.

### Time vs. Service



*Figure 5: The service duration in milliseconds versus the service number for transform frame-only write-back with a 35-millisecond deadline.*



## **Appendix A: Code**

Note: code was not included here for this assignment due to the large size of the files.